



AI in Games

Overview

- A smattering of game AI techniques
 - Goal-based planning
 - Influence maps
 - Neural networks
 - Evolved virtual creatures
 - NERO
 - Learning bots in Quake
 - Modeling randomness

Classical Game AI

- Finite state machine
 - Hard-coded behaviors, specify all possible combinations of actions
- A-star path planning
 - On top of a ~~waypoint graph~~ navigation mesh
 - Also local steering forces, object avoidance, etc.
- These techniques are widely used today
...But are not the focus of today's lecture

Alternative: Goal-Driven AI

- Hierarchical set of goals
 - e.g. "buy sword" => "obtain gold" + "go to smithy"
- Plan & perform actions to achieve goals
 - Planner finds a sequence of actions that will achieve a goal
 - Sequence is *not* hard-coded, discovered on the fly
- May also use A-star path planning

Planning Challenges

- Combinatorial explosion of world states
 - Need to simplify representation of world
 - Only model what is relevant to a specific agent
- Efficiently searching through possible actions
 - Can model as graph search problem
- Biasing towards certain actions
 - Weights on graph

Case Study: F.E.A.R

- 2005 horror FPS
- Ranked #2 Most influential AI Game

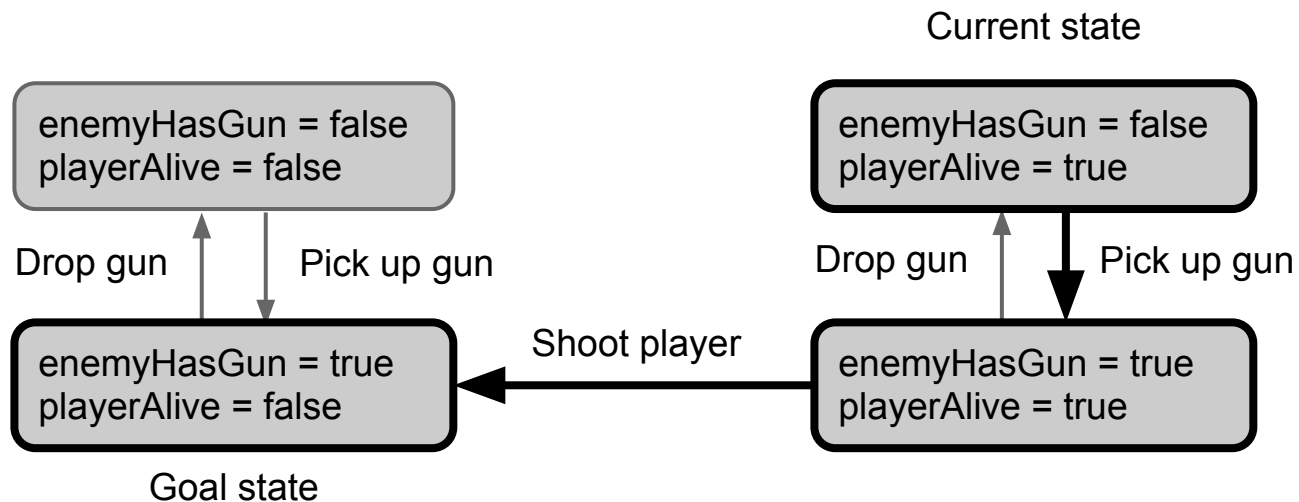


F.E.A.R. AI

- Uses GOAP (Goal-Oriented Action Planning)
 - Consists of goals and actions
 - Goal: certain state of the world we want to reach
 - Action: set of preconditions and effects
- Solved using graph search
 - Nodes are world states and edges are actions
 - Search from current state toward goal
 - Edges are directional (directed graph)
 - Analogy to chess AI: nodes are board states and edges are moves

GOAP Example

- World state:
 - bool enemyHasGun
 - bool playerAlive
- Only models states relevant to agent



GOAP

- Different from FSMs
 - FSM: Spend all time in nodes (actions), transitions are instant
 - GOAP: Spend all time in transitions (actions), nodes are instant
- Advantages over FSMs
 - Don't need to encode all possible transitions from each action to each other action
 - Simpler to specify and easier to scale
 - Just specify preconditions and effects for each action

GOAP in F.E.A.R.


- Graph is dynamic
 - Edges are actions with preconditions and effects
 - Edges come and go based on game state
- Procedural preconditions
 - Taking an action may be currently impossible
 - Example: Escape through door requires door unlocked
- Procedural effects
 - Effects take time to execute and may fail
 - Example: Firing a projectile is blocked by something

GOAP in F.E.A.R.

- Don't want all agents to act the same
 - Vary edge weights based on individual preference (aggressive vs careful individual)
 - Need variety in available actions (run, crouch, dodge, roll, slide)
- Hard for player to understand AI's behavior
 - Add audio cues for player
 - "Send in reinforcements!"

Influence Maps

- Goal: decide which team controls an area
 - Used as an input to higher level decision algorithms



Influence Maps

- Details

- Numeric value that varies throughout the world
- Positive for team A, negative for team B
- Absolute value indicates strength of influence
- Can also encode other information: safety, congestion, etc...

- Advantages

- Entire influence map not updated immediately, will "remember" recent history
- Helps reason about strategy in complex worlds (open areas, choke points)
- Commonly used in RTS games

Influence Maps

- Algorithm

- Influence map starts off at zero
- Sources of influence are marked
 - Entities (friends and foes)
 - Events like grenade explosions, taking damage
- Influence is propagated (diffused) to neighbors


- Parameters

- Momentum (m): How fast do new values overwrite old values?
- Decay (d): How fast do values decrease from the source?

$$x(t+1) = (1-m) \cdot x(t) + m \cdot \sum_{a \in agents} (x_a \cdot e^{-dist_a \cdot d})$$

Influence Maps

- Influence maps can be combined with graphs



Neural Networks



Not used in commercial games

Evolved Virtual Creatures

- Develop a creature and a controller using a genetic algorithm
 - Research project by Karl Sims in 1994
- Start with randomized creatures
 - Configuration of body parts in a tree
 - Brain is a graph
 - Sensors (joint angle, contact, light)
 - Neurons (sum, sin, sigmoid, ...)
 - Effectors (a degree of freedom)
 - Goal defines fitness function
 - Example: jumping measures height of lowest body part

Evolved Virtual Creatures

- Evolve the current generation by randomly applying one of:
 - Mutation (asexual reproduction)
 - Crossovers (sexual reproduction)
 - Grafting (sexual reproduction)



Evolved Virtual Creatures



http://www.youtube.com/watch?v=JBgG_VSP7f8

Genetic Algorithms in Games


- Hardly ever used
 - Developers want fine control over behavior, often just enumerate all the cases
 - Difficult to test
 - Genetic algorithms take thousands of iterations to converge
- Creating faster genetic algorithms is an active area of research

Case Study: NERO

- Full name: Neuro-Evolving Robot Operatives
- Machine-learning based game
 - Started in 2003 as research project at UT Austin
- RTS-style gameplay with two phases of play
 - Phase 1: Train agents in sandbox scenarios
 - Phase 2: Simulate a battle between two groups of agents
- Works through a real-time genetic algorithm
 - NEAT (Neuro-Evaluation of Augmented Topologies)

NEAT Algorithm

- Neural network evolves over time
 - Evolution can change both connection weights and network structure



<http://www.youtube.com/watch?v=T4EopjWkLtl>

Real-time NEAT

- Start with minimally-connected neural network
 - Add connections that help solve the problem
- Continuously adapts a small group of agents
 - Doesn't use generations: this takes too long
 - Each "brain" has a evaluation time, when it runs out, it's replaced by merging two high-performance brains
 - Removed "brain" is shelved for later evaluation
- Open-source
 - <http://nn.cs.utexas.edu/keyword?rtneat>

Real-time NEAT in NERO

Every n game ticks:

1. Remove agent with worst adjusted fitness
2. Re-estimate fitness for all species
3. Choose a parent species
4. Adjust species boundaries and reassign all agents to species
5. Add new agent to the world

Case Study: Learning Bots in Quake

- Remco Bonse et. al. in 2004
- Use the Q-learning algorithm to train a neural network for an AI bot in Quake III
 - One-on-one game against preprogrammed AI
 - Replaced combat movement subsystem with NN
 - Reward is given for avoiding damage
- State:
 - Distance and angle of opponent
 - Distance and angle of nearest projectile
- Actions:
 - All 18 combinations of WASD + jump

Case Study: Learning Bots in Quake

- Q-learning algorithm:
 - Agent is in some state $s \in S$
 - Agent can take some action $a \in A$
 - State/action quality function: $Q : S \times A \rightarrow \mathbb{R}$
- Agent gets reward for each state change
 - Make a correction to Q based on new information
 - Learning rate: 0 = learn nothing, 1 = ignore past
 - Discount factor: 0 = short-term, 1 = long-term

$$Q(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha_t(s_t, a_t)}_{\text{learning rate}} \times \left[\underbrace{R_{t+1}}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \underbrace{\max_{a_{t+1}} Q(s_{t+1}, a_{t+1})}_{\text{max future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right]$$

Case Study: Learning Bots in Quake

- Conclusion

- Best parameters were to play to 100,000 frags on a network of 15 neurons (10 hours training time)
- Learned bot wins twice as often against a preprogrammed bot
- Human players didn't notice a difference
- Limited interaction time in a FPS

Randomness and AI

- One motivation for AI is variability
- Adding randomness can help
 - Easy way of varying AI behaviors
 - Need to be careful about how user's perceive randomness


Randomness and Human Perception




Image from dilbert.com

Randomness and Human Perception

- Which is more random?



or




Randomness and Human Perception

- Humans are illogical
 - Conditioned to see patterns
 - Patterns are more memorable
 - "He got 3 but I only got 1"
- Gambler's Fallacy
 - We think outcomes depend on past events for independent events with fixed probabilities
 - Not just gamblers, everyone does this
 - "His luck will run out sooner or later"


Law of Large Numbers

- Average approaches the expected value as the number of trials increases



Law of Large Numbers

- Average approaches the expected value as the number of trials increases



Law of Small Numbers?

- Make randomness easier to understand
 - Random choices without replacement
 - Keep track of past choices and prevent them from being chosen again too soon
 - Humans expect random binary choices will alternate up to 40% more than is mathematically reasonable¹
 - Example: If a player loses a game involving randomness 4 times in a row, is that a good experience?

¹ A Meta-Analysis of Randomness in Human Behavioral Research

C++ Tip of the Week

- Multiple inheritance
 - C++ classes can have multiple superclasses
 - Can be used to emulate Java interfaces

```
// Declare interfaces using pure virtual functions
struct Reader { virtual char read() = 0; };
struct Writer { virtual void write(char) = 0; };

// File * can be cast to Reader * and Writer *
struct File : Reader, Writer {
    char read() { ... }
    void write(char) { ... }
};
```

C++ Tip of the Week

- More powerful than Java interfaces
 - All superclasses can have state

```
struct Timestamp {  
    const int creationTime;  
    Timestamp() : creationTime(time(0)) {}  
};  
  
struct UniqueID {  
    const int id; static int nextID;  
    UniqueID() : id(nextID++) {}  
};
```


```
struct Entity : Timestamp, UniqueID {};
```

C++ Tip of the Week

- The diamond problem and virtual inheritance

```
struct Base { int num; };  
struct Derived1 : Base {};  
struct Derived2 : Base {};  
struct Join : Derived1, Derived2 {};
```

```
Join join;  
cout << join.num; // error, join has two copies of num  
cout << join.Derived1::num;  
cout << join.Derived2::num;
```




C++ Tip of the Week

- The diamond problem and virtual inheritance

```
struct Base { int num; };  
struct Derived1 : virtual Base {};  
struct Derived2 : virtual Base {};  
struct Join : Derived1, Derived2 {};
```

```
Join join;
```

```
cout << join.num; // works, join has one copy of num
```



C++ Tip of the Week

- Tricky case

- Multiple pointers for join (base1 != base2)
- Using virtual inheritance here would cause an error

```
struct Base { virtual void foo() = 0; };  
struct Derived1 : Base { void foo() { cout << 1; } };  
struct Derived2 : Base { void foo() { cout << 2; } };  
struct Join : Derived1, Derived2 {};
```

```
Join join;
```

```
Base *base = &join; // error, ambiguous conversion
```

```
Base *base1 = static_cast<Derived1 *>(&join);
```

```
Base *base2 = static_cast<Derived2 *>(&join);
```

C++ Tip of the Week

- Tricky case

- Fix: define the behavior we want by overriding foo() again in the class with multiple inheritance

```
struct Base { virtual void foo() = 0; };  
struct Derived1 : virtual Base { void foo() { ... } };  
struct Derived2 : virtual Base { void foo() { ... } };  
struct Join : Derived1, Derived2 {  
    void foo() { Derived1::foo(); Derived2::foo(); }  
};
```

```
Join join;
```

```
Base *base = &join; // works, only one foo() implementation
```


C++ Tip of the Week

- Common use case: policy-based designs
 - Compile-time version of the strategy pattern
 - Each template parameter provides part of behavior
 - Compiler can inline and optimize policy methods

```
template <typename K, typename V, class Alloc, class Hash>
struct HashMap : Alloc, Hash {
    V &get(const K &key) {
        Bin &bin = bins[hash(key)];
        if (!bin.contains(key)) bin.insert(key, alloc());
        return bin.find(key);
    }
};

HashMap<Point, Entity, MemoryPool, SecureHash> data;
```

C++ Tip of the Week

- C++ mixins

- Alternative to multiple inheritance
- Inherit from a template parameter
- Use initializer functions instead of constructors if you need arguments on construction

```
template <typename Base>  
struct Timestamp : Base { ... };
```

```
template <typename Base>  
struct UniqueID : Base { ... };
```

```
Timestamp<UniqueID<Entity>>> entity;
```

Final Project Questions?

References

- Using Randomness in AI: Both Sides of the Coin (Dave Mark and Brian Schwab)
- http://web.media.mit.edu/~jorkin/gdc2006_orkin_jeff_fear.pdf
- Evolving Virtual Creatures (Karl Sims)
- Evolving Neural Network Agents in the NERO Video Game (Kenneth Stanley et al.)
- A Meta-Analysis of Randomness in Human Behavioral Research (Summer Ann Armstrong)
- Learning Agents in Quake III